

James C. Davis

davisjam@purdue.edu

<https://davisjam.github.io/>

RESEARCH INTERESTS

My research focuses on **software quality**. I use qualitative and quantitative methods to identify socio-technical defects and shortcomings. I identify and resolve problems using systems and security principles. For example:

- **Troublesome tools:** I have given a strong empirical foundation to the study of regular expressions (FSE'18, FSE'19, ASE'19a, ASE'19b).
- **Emerging paradigms:** I have studied the systems and security issues that arise in event-driven programming as embodied in the Node.js framework (EuroSys'17, USENIX Security'18). I have also examined security issues in the use of the query language GraphQL (ICSOC'19, FSE'20).
- **Distributed systems:** I hold several patents on testing distributed storage systems, and have researched performance issues in stream processing engines (USENIX ATC'19).

EDUCATION

Ph.D Computer Science

2015-2020 (expected)

Virginia Tech

Advisor: Dr. Dongyoon Lee

Dissertation: *On the Impact and Defeat of Regular Expression Denial of Service*

B.S. Computer Science, B.S. Mathematics

2008-2012

Clarkson University

Research mentor: Dr. Takashi Nishikawa

Honors Thesis: *Relating Synchronizability to the Topological Properties of Networks Using a Linear Classifier*

EMPLOYMENT

Assistant Professor, Electrical and Computer Engineering

Fall 2020-present

Purdue University

Graduate Research Assistant

2016-2020

Virginia Tech

- **Practice-motivated systems design:** Re-designing regex engines based on practitioners' perspectives.
- **Security in emerging paradigms:** Denial of service in event-driven frameworks and GraphQL.

Intern, Microsoft Research (RiSE group: Cloud Security)

Summer 2019

Microsoft Research, Redmond, WA

Project sponsor: Dr. Patrice Godefroid

- Techniques and tools for web API security testing.

Intern, IBM Research (Storage)

Summer 2018

IBM Research, Almaden, CA

Project sponsor: Dr. Deepavali Bhagwat

- Provenance collection system for machine learning applications (VLDB'20).

Software Engineer, IBM (GPFS)

2012-2015, Summer 2016, Summer 2017

IBM, Poughkeepsie, NY

- Developed distributed applications and tooling for cluster management and file system testing.
- Worked with and trained test teams in the US, the UK, Mexico, Germany, India, and China.
- Created five patents on techniques for file system testing.

CONFERENCE PAPERS

Cha, Wittern, Baudart, **Davis**, Mandel, Laredo. *A Principled Approach to GraphQL Query Cost Analysis*. Proceedings of the 28th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (**ESEC/FSE'20**). *ACM Distinguished Paper Award*.

Rupprecht, **Davis**, Arnold, Gur, Bhagwat. *Improving Reproducibility of Data Science Pipelines through Transparent Provenance Capture*. Proceedings of the 46th International Conference on Very Large Data Bases (**VLDB'20 Industry track**).

Davis, Moyer, Kazerouni, and Lee. *Testing Regex Generalizability And Its Implications: A Large-Scale Many-Language Measurement Study*. Proceedings of the 34th IEEE/ACM International Conference on Automated Software Engineering (**ASE'19**).

Michael, Donohue, **Davis**, Lee, and Servant. *Regexes are Hard: Decision-making, Difficulties, and Risks in Programming Regular Expressions*. Proceedings of the 34th IEEE/ACM International Conference on Automated Software Engineering (**ASE'19**). *ACM Distinguished Paper Award*.

Davis, Michael, Coghlan, Servant, and Lee. *Are Regular Expressions a Lingua Franca? An Empirical Study on the Re-use and Portability of Regular Expressions*. Proceedings of the 27th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (**ESEC/FSE'19**).

Wittern, Cha, **Davis**, Baudart, Mandel. *An Empirical Study of GraphQL Schemas*. Proceedings of the 17th International Conference on Service-Oriented Computing (**ICSOC'19**).

Fu, Ghaffar, **Davis**, and Lee. *EdgeWise: A Better Stream Processing Engine for the Edge*. 2019 USENIX Annual Technical Conference (**USENIX ATC'19**).

Davis, Coghlan, Servant, and Lee. *The Impact of Regular Expression Denial of Service (REDOS) in Practice: an Empirical Study at the Ecosystem Scale*. Proceedings of the 26th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (**ESEC/FSE'18**). *ACM Distinguished Paper Award*.

Davis, Williamson, and Lee. *A Sense of Time for JavaScript and Node.js: First-Class Timeouts as a Cure for Event Handler Poisoning*. Proceedings of the 27th USENIX Security Symposium (**USENIX Security'18**).

Davis, Thekumparampil, and Lee. *Node.fz: Fuzzing the Server-Side Event-Driven Architecture*. Proceedings of the Twelfth European Conference on Computer Systems (**EuroSys'17**).

JOURNAL ARTICLES

Ozkan, Davis, **Davis**, James, Murzi, Knight. *Expectations and Experiences of Short-Term Study Abroad Leadership Teams*. Journal of International Engineering Education (**JIEE**).

SHORT PAPERS

Davis. *On the Impact and Defeat of Regex DoS*. **ACM SRC – Grand Finals. Second place, graduate student division**.

Rupprecht, **Davis**, Arnold, Lubbock, Tyson, and Bhagwat. *Ursprung: Provenance for Large-Scale Analytics Environments*. Proceedings of the 2019 International Conference on Management of Data (**SIGMOD'19 Demo**).

Davis, Kildow, and Lee. *The Case of the Poisoned Event Handler: Weaknesses in the Node.js Event-Driven Architecture*. Proceedings of the 10th European Workshop on Systems Security (**EuroSec'17**).

PATENTS

Davis, **Davis**. *Injection of Simulated Hardware Failure(s) in a File System for Establishing File System Tolerance-to-Storage-Failure(s)*. IBM, U.S. patent application 20200264961. U.S. patent pending.

Davis, **Davis**. *Verification of the integrity of data files stored in copy-on-write (CoW) based file system snapshots*. IBM, U.S. patent application 20200242075. U.S. patent pending.

Davis, Davis. *File metadata verification in a distributed file system*. IBM, U.S. patent 10,678,755B2, granted Jun. 9, 2020..

Davis, **Davis**. *Testing of lock managers in computing environments*. IBM, U.S. patent 10,061,777 B1, granted Aug. 28, 2018.

Davis, Davis, Knop. *Detection of file corruption in a distributed file system*. IBM, U.S. patent 10,025,788, granted Jul. 17, 2018.

TEACHING

Instructor, ECE 368 Data Structures Fall 2020
Purdue University

Co-advisor, VIP: Open-Source TensorFlow Software Fall 2020
Purdue University

Instructor, Data Structures and Algorithms Fall 2019
Virginia Tech

Instructor, Introduction to Programming in Python Spring 2019
Virginia Tech

Track Leader, Rising Sophomore Abroad Program Spring 2018, Spring 2019
Virginia Tech

STUDENTS

Undergraduates

Albert Sun (BS, Purdue University)	2020-present
Jitesh Motati (BS, Purdue University)	2020-present
Tuhin Sarkar (BS, Purdue University)	2020-present
Jenna Ryan (BS, Clarkson University)	2018-present
Jonathan Alexander (BS, Virginia Tech)	2018-2019
Christy Coghlan (BS, Virginia Tech)	2018
Celine Stewart (BS, Virginia Tech)	2017-2019

INVITED TALKS

Regexes Awry: Characterizing and Defeating Regex-based Denial of Service <i>Clarkson University CS department colloquium</i>	2020
Regexes are Hard: Qualitative and Quantitative Perspectives <i>NC State CS department colloquium</i>	2019
The Dangers of Copy/Pasting Code <i>Episode of the Podcast “The Secure Developer”: https://tinyurl.com/DavisResearchPodcast</i>	2019
Regexes in the Wild <i>Virginia Tech department seminar</i>	2019
Academic Perspectives on Node.js <i>Node.js Collaborator Summit, Vancouver</i>	2018
International Engineering <i>Rising Sophomore Abroad Program, Virginia Tech</i>	Annual, 2015-2019

OTHER NOTABLE ACTIVITY

Disclosed DoS vectors in Python core and Node.js core <i>Python: CVE-2018-1060, CVE-2018-1061; Node.js: CVE-2018-7158</i>	
Guide: “Don’t Block the Event Loop (or the Worker Pool)” <i>https://nodejs.org/en/docs/guides/dont-block-the-event-loop/</i>	
VT Intramural Racquetball Champion (Singles, Doubles)	Spring, Fall

AWARDS AND RECOGNITION

ACM Distinguished Paper Award, ESEC/FSE 2020	2020
Outstanding Graduate Student Service Award, CS@VT	2020
Second place, Grand Finals of the ACM Graduate Student Research Competition	2020
First place, Graduate Student Research Competition, ESEC/FSE 2019	2019
ACM Distinguished Paper Award, ASE 2019	2019
ACM Distinguished Paper Award, ESEC/FSE 2018	2018
Microsoft Security Researcher Acknowledgments (Regex DoS)	2018
Graduate Fellow, VT Academy for Global Engineering	2019-2020
IBM Significant Contributor Award (Node.js)	2018
IBM Poughkeepsie’s New hire of the month	2014
Frederica Clarkson Award	2012
Clarkson University’s Outstanding Senior (x2): Mathematics, Computer science	2012
Clarkson University Phalanx Commendable Leadership	2011

EXTERNAL SERVICE

Reviewer, ACM Transactions on Software Engineering (TSE)	2020-present
Member, ICSE Demonstrations Track	ICSE 2021
Member, ESEC/FSE Artifact Evaluation Committee	ESEC/FSE 2020
Member, CGO Artifact Evaluation Committee	CGO 2019
Sub-reviewer: ASPLOS'18, EuroSys'18, MASCOTS'18, HPCA'19, CGO'19, ISMM'19	2016-2019
Regional Judge, ACM ICPC	Fall 2015

INTERNAL SERVICE

Committee member, Purdue ECE Undergraduate Curriculum Committee	2020-present
President, VT CS Graduate Student Council	2018-2019
Organizer, VT Systems Reading Group	Spring 2017-2020

SELECTED OPEN-SOURCE PROJECTS

safe-regex	Check if your regex is super-linear. 4.7M dependents.
regexp-tree	Analysis tools for regular expressions.
marked	Regex-based Markdown parser (I consult on ReDoS). 354K dependents.

PROFESSIONAL MEMBERSHIPS

Member, Association for Computing Machinery
